

Introduction to Neural Network, Convolution Neural Networks, and Deep Learning

Paul Rodriguez, PhD
(SDSC)

08/10/2023

Schedule overview

Mai Nguyen, Paul Rodriguez

- **Yesterday - Scaling**

- R on HPC
- Spark

- **Today – Deep Learning**

- Part 1A Intro to NN/CNN/Deep Learning
- Part 1B Practical Guidelines and Multinode execution
- DL Layers and Models
- DL Transfer Learning
- DL Functional API, Special Connections, Transformers

Outline

- **Part I A**

Overview of Neural Networks (aka Multilayer Perceptron)

Convolution Neural Networks and Scaling

Exercise, MNIST classification

- **Part I B**


Practical Guidelines: Hyperparameters, Workflows, GPUs

Exercise, Multinode MNIST

Logistic Regression to Neural Network


$$f(x, b, w) = \frac{\exp(b+wx)}{1 + \exp(b+wx)} = \frac{1}{1 + \exp(-(b+wx))}$$

for parameters: $b = 0$, $w_1 = 1$







Logistic Regression to Neural Network


$$f(x, b, w) = \frac{\exp(b + w * x)}{1 + \exp(b + w * x)}$$




Logistic function w/various weights



Example: 1 input into 2 logistic units




Example: 1 input into 2 logistic units with these activations





Example: 1 input into 2 logistic units with these activations

If you add these 2 units into a final output unit what would the output function look like?



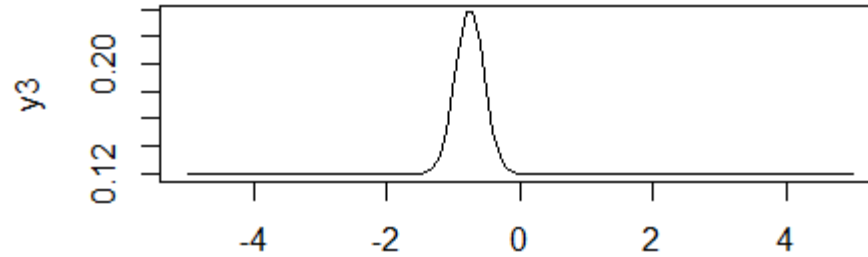
Example: 1 input into 2 logistic units with these activations

If you add these 2 units into a final output unit what would the output function look like?




Higher level function combinations


```
x=seq(-5,5,.1)
y1=1/(1+exp(10+ 10*x))
y2=1/(1+exp(-5+(-10)*x))
y3=1/(1+exp(1+1*y1+1*y2))
plot(x,y3,type="l")
```



```
y4=1/(1+exp(10+ (-10)*x))
y5=1/(1+exp(-5+(10)*x))
y6=1/(1+exp(1-1*y4-1*y5))
plot(x,y6,type="l")
```

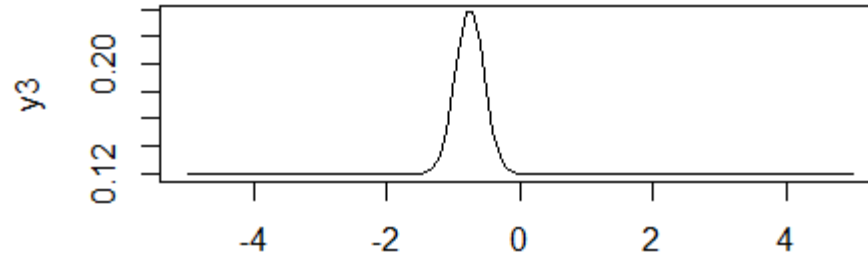


```
y7=1/(1+exp(1+2*y3+1*y6))
plot(x,y7,type="l")
```




Higher level function combinations

```
x=seq(-5,5,.1)
y1=1/(1+exp(10+ 10*x))
y2=1/(1+exp(-5+(-10)*x))
y3=1/(1+exp(1+1*y1+1*y2))
plot(x,y3,type="l")
```




Multiple layer networks can represent any logical or real-valued functions (unbiased, but potential to overfit)

```
y4=1/(1+exp(10+ (-10)*x))
y5=1/(1+exp(-5+(10)*x))
y6=1/(1+exp(1-1*y4-1*y5))
plot(x,y6,type="l")
```



```
y7=1/(1+exp(1+2*y3+1*y6))
plot(x,y7,type="l")
```





Logistic to Neural Network model

$$f(x, b, w) = \frac{\exp^{(b+wx)}}{1 + \exp^{(b+wx)}}$$

Draw out function as a little graph, 1 input


Logistic to Neural Network model


$$f(x, b, w) = \frac{\exp^{(b+w*x)}}{1 + \exp^{(b+wx)}}$$




Draw out function as a little graph, 1 input

Logistic to Neural Network model


$$f(x, b, w) = \frac{\exp^{(b+w*x)}}{1 + \exp^{(b+wx)}}$$




Draw out function as a little graph, 1 input

logistic function will transform input to output – call it the ‘activation’ function


Using 2 input units, the graph model would be:



We usually don't draw the bias.


We assume inputs*weights are summed (a dot product)

Using 2 input units, 2 intermediate units, and 1 output:




Call these "hidden units"

Using 2 input units, 2 intermediate units, and 1 output:




For X a $P \times 1$ vector, we can set up indices in a weight matrix so that: $W \cdot X =$ incoming activations from previous layer

More generally, we can add a hidden layer,
and have many inputs and outputs




A “Multilayer Perceptron”


1 Input layer 1 Hidden layer 1 Output




A “Multilayer Perceptron”




A “Multilayer Perceptron”



A “Multilayer Perceptron”




Algorithm steps



1. FORWARD PROPAGATE
AN ENTIRE BATCH OF
INPUTS

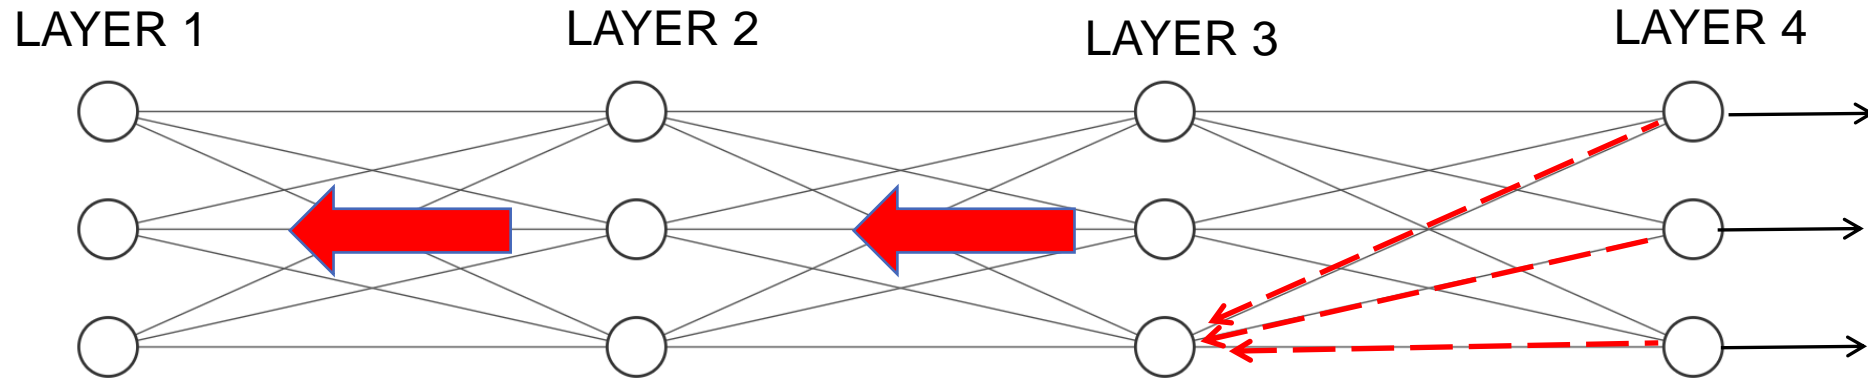
Algorithm steps



2. BACKWARD PROPAGATE ERROR FOR WHOLE BATCH USING DERIVATIVE CHAIN RULE:

$$\frac{dE}{dw_{mp}} = \sum_k^K \frac{dE_k}{d\hat{y}_k} \frac{d\hat{y}_k}{da_k} \frac{da_k}{dh_m} \frac{dh_m}{da_m} \frac{da_m}{dw_{mp}}$$

Algorithm steps and Vanishing Gradients

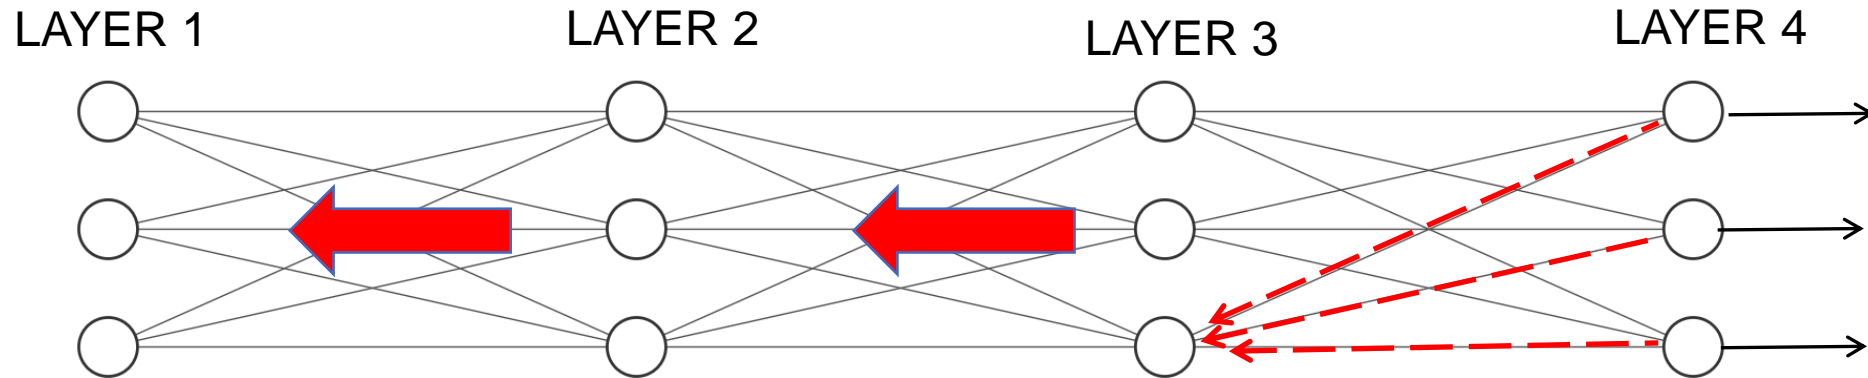


2. BACKWARD PROPAGATE ERROR FOR WHOLE BATCH USING DERIVATIVE CHAIN RULE:

Note: As you go farther back, the error information gets diluted and the error gradient starts 'vanishing'

$$\frac{dE}{dw_{mp}} = \sum_k^K \frac{dE_k}{d\hat{y}_k} \frac{d\hat{y}_k}{da_k} \frac{da_k}{dh_m} \frac{dh_m}{da_m} \frac{da_m}{dw_{mp}}$$

Algorithm steps and Vanishing Gradients



2. BACKWARD PROPAGATE ERROR FOR WHOLE BATCH USING DERIVATIVE CHAIN RULE:

Note: As you go farther back, the error information gets diluted and the error gradient starts ‘vanishing’

A different activation function helps ...

$$\frac{dE}{dw_{mp}} = \sum_k^K \frac{dE_k}{d\hat{y}_k} \frac{d\hat{y}_k}{da_k} \frac{da_k}{dh_m} \frac{dh_m}{da_m} \frac{da_m}{dw_{mp}}$$


The rectified linear unit (RELU)

RELU (rectified linear unit)

RELU activation function

It is unscaled (bad!)

But df/da is constant (good!)



$$f(a) = \begin{cases} a & a > 0 \\ 0 & a \leq 0 \end{cases}$$

where $a = XW$

Overall, RELU mitigates vanishing gradients

The Neural Network Algorithm

INITIALIZE weights to small value (for example: $\pm < 0.3$)

LOOP until stopping criterion:

The Neural Network Algorithm

INITIALIZE weights to small value (for example: +/- <0.3)

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all error derivatives to *minimize Loss*

The Neural Network Algorithm

INITIALIZE weights to small value (for example: +/- <0.3)

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all error derivatives to *minimize Loss*

UPDATE WEIGHTS: $w \leftarrow w - \text{learning_rate} * \frac{dL}{dw}$

The Neural Network Algorithm

INITIALIZE weights to small value (for example: +/- <0.3)

LOOP until stopping criterion:

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all error derivatives to *minimize Loss*

UPDATE WEIGHTS: $w \leftarrow w - \text{learning_rate} * \frac{dL}{dw}$

STOP: when validation error reaches minimum or after a max number of epochs

The Neural Network Algorithm [and heuristics]

INITIALIZE weights to small value (for example: +/- <0.3)

LOOP until stopping criterion:

[work in batches of input]

FORWARD PROPAGATION: calculate all node activations

BACKWARD PROPAGATION: calculate all error derivatives to *minimize Loss*

UPDATE WEIGHTS: $w \leftarrow w - \text{learning_rate} * \frac{dL}{dw}$

[adapt learning rate,
use momentum]

STOP: when validation error reaches minimum or after a max number of epochs

[several metrics of loss are possible]

Neural Network main options to choose:

- 1 Architecture: number of hidden units & layers
- 2 Optimizer and learning rate
- 3 Loss function depends on task
- Note: more hidden layers, more hidden units => more potential for overfitting

terminology and cheat sheet on output activations (for reference):

Type of Problem	Y outputs	Output Activation Function (this gives a SCORE \hat{Y} :)	Output PREDICTION (what you decide to predict)	Output Loss Function	Evaluative Measure
Regression: map into to K real valued predictions	if $Y \in (-\infty, +\infty)^K$	$\hat{Y} = XW$	\hat{Y} :	Sum Squared Error (SSE)	Mean Squared Error (MSE)
Multivariate output of 0's and 1's	if $Y \in [0, 1]^K$	$\hat{Y} = \frac{1}{1 + \exp^{-(XW)}}$	1 or 0	SSE	MSE
Binary Classification	if $Y \in \{0, 1\}$	$\hat{Y} = \frac{1}{1 + \exp^{-(XW)}}$	A probability given by \hat{Y} : $P(y = 1 x)$	Cross Entropy $L = -y \log(\hat{y}) - (1 - y)(\log(\hat{y}))$	Accuracy, ROC
Multiclassification	if $Y \in \{0, 1\}^K$	$\hat{Y}_k = \frac{\exp^{-(XW_k)}}{\sum_k \exp^{-(XW_k)}}$	Max class	Cross Entropy $L = - \sum_k y_k \log(\hat{y}_k)$	Accuracy

Summary:

Pro:

Multilayer Perceptron, and Neural Networks in general, are flexible powerful learners

Hidden layers learn a nonlinear transformation of input

Summary:

Pro:

Multilayer Perceptron, and Neural Networks in general, are flexible powerful learners

Hidden layers learn a nonlinear transformation of input

Con:

Lots of parameters

Hard to interpret

Needs more data

**A neural network can discover visual features using
'convolutions'**


Next: Image classification of digits



Image features

- **MNIST - A database of handwritten printed digits**
(National Inst. of Standards and Technology)





How to classify digits?





Let's zoom into 5x6 window of pixels near the tip of '7'

Take a 3x3 patch of pixels and apply a 'filter' template – designed to find an edge




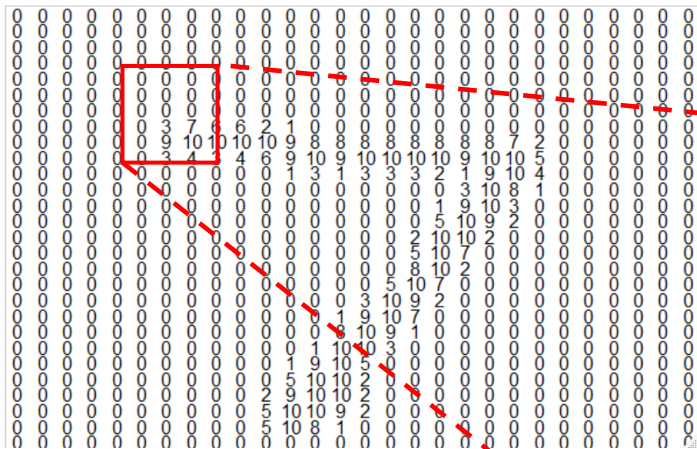
X

-1	0	+1
-1	0	+1
-1	0	+1

Let's zoom into 5x6 window of pixels near the tip of '7'

Take a 3x3 patch of pixels and apply a 'filter' template – designed to find an edge

1. Multiply 3x3 patch of pixels with 3x3 filter



(our weight parameters)



X

-1	0	+1
-1	0	+1
-1	0	+1

1. Multiply 3x3 patch of pixels with 3x3 filter “W”

Let's zoom into 5x6 window of pixels near the tip of '7'


Take a 3x3 patch of pixels and apply a 'filter' template – designed to find an edge



-1	0	+1
-1	0	+1
-1	0	+1


X

1. Multiply 3x3 patch of pixels with 3x3 filter "W"





2. Put sum in new cell of output map

3. Slide filter horizontally to get next output value



3			

3. Slide filter horizontally to get next output value



X


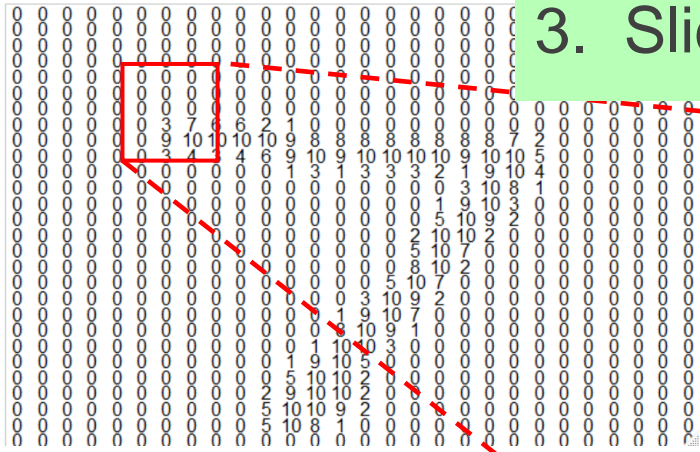
-1	0	+1
-1	0	+1
-1	0	+1

1. Multiply 3x3 patch of pixels with 3x3 filter "W"

2. Put sum in new cell of output map

3	7		

3. Slide filter horizontally to get next output value



-1	0	+1
-1	0	+1
-1	0	+1


X

1. Multiply 3x3 patch of pixels with 3x3 filter “W”

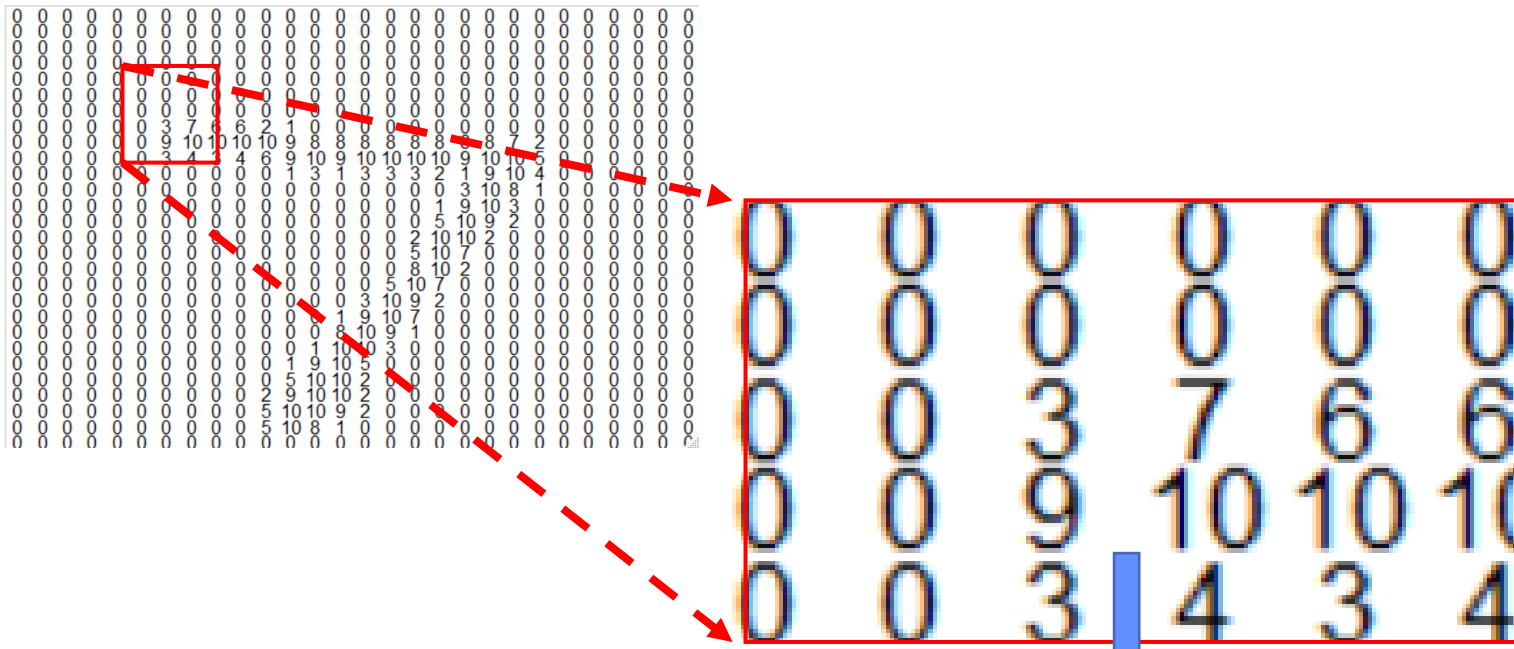
NOTE: sliding a filter is known as a “convolution” operation

3	7	3	

2. Put sum in new cell of output map



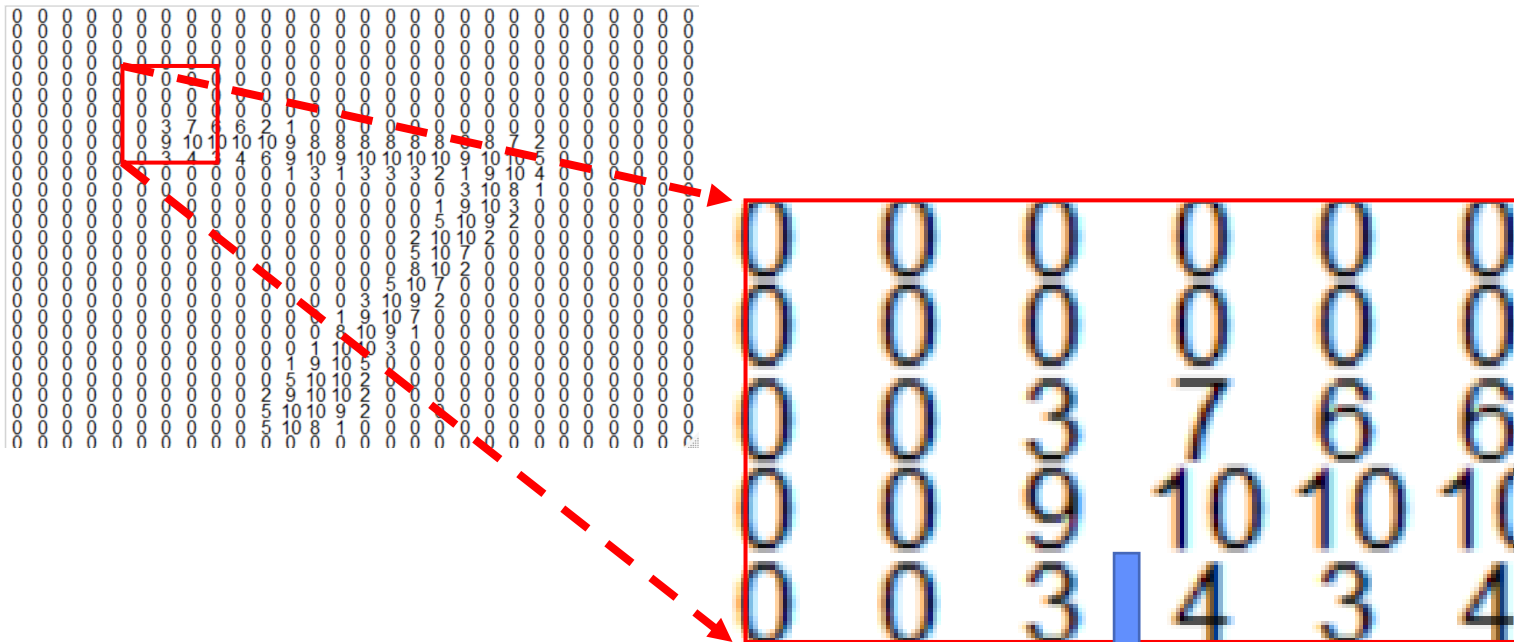
After vertical and horizontal sliding the 5x6 patch is now a 3x5 feature map.



3	7	3	-1
12	17	4	-1
15	21	4	-1

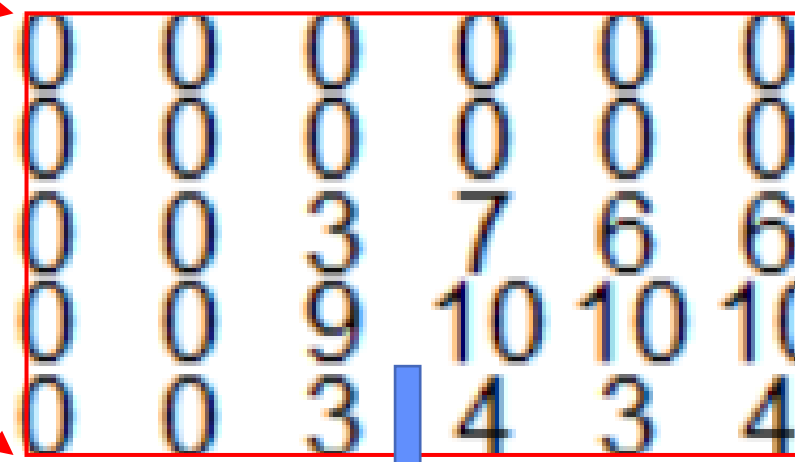
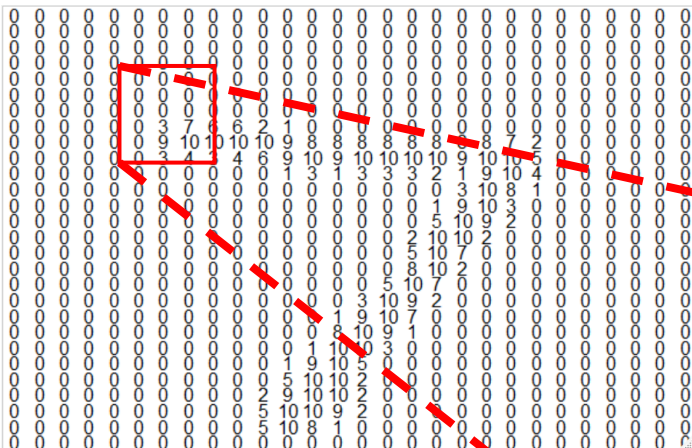
After vertical and horizontal sliding the 5x6 patch is now a 3x5 feature map.

What do the highest values in the feature map represent?



3	7	3	-1
12	17	4	-1
15	21	4	-1

Optional next step:
Use another filter, and take maximum over elements -
“max pooling”

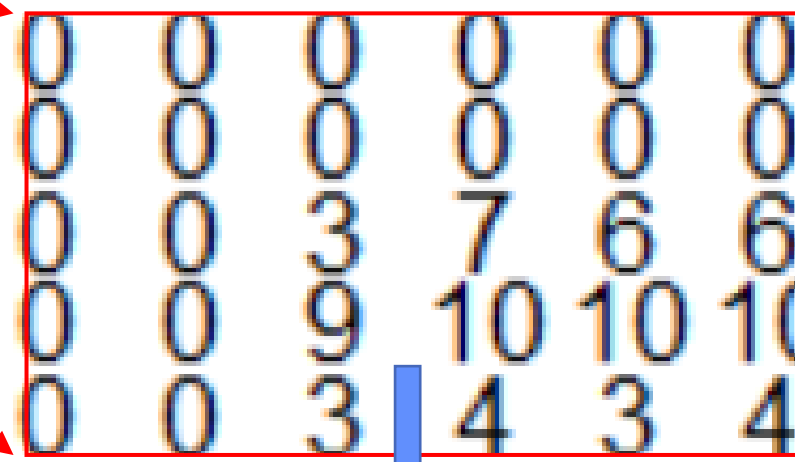
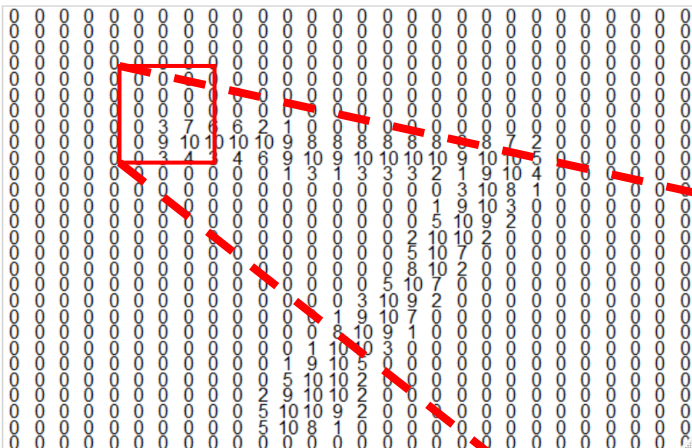


3	7	3	-1
12	17	4	1
15	21	4	-1

Optional next step:
Use another filter, and take maximum over elements -
“max pooling”

2x2 filter has max=17

17		

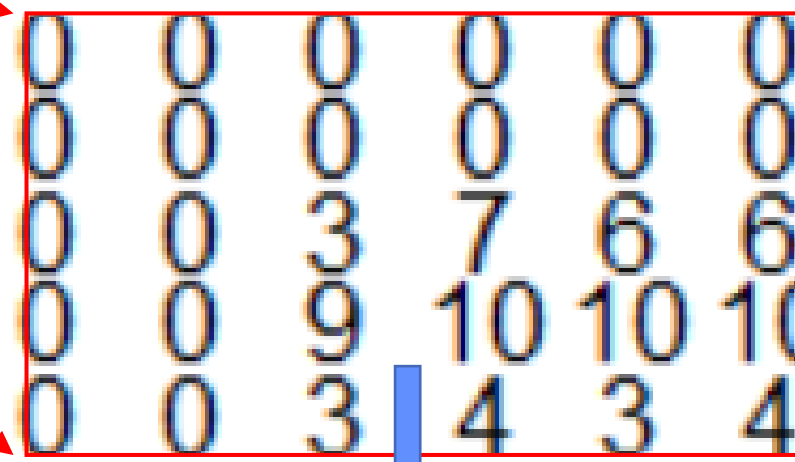
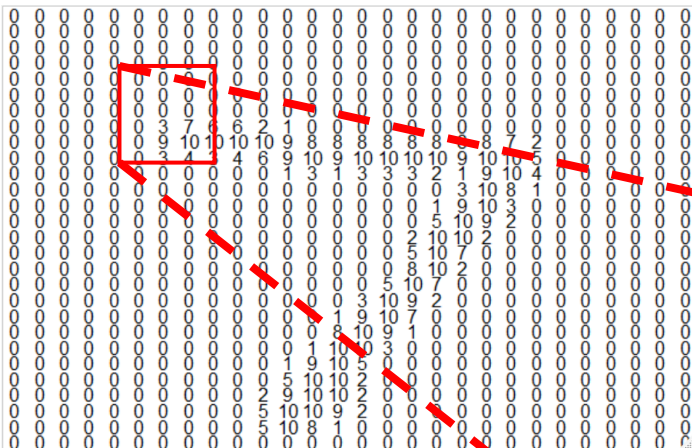


3	7	3	-1
12	17	4	-1
15	21	4	-1

Optional next step:
Use another filter, and take maximum over elements -
“max pooling”

Slide filter ...

17	17	4
21	21	4





3	7	3	-1
12	17	4	-1
15	21	4	-1

After convolution and pooling, the 5x6 patch is **transformed** into a 2x3 feature map of 'edge gradients'

Slide filter ...

17	17	4
21	21	4



A convolution of one filter is applied to the entire image across and down.

*The entire 28x28 input is **transformed** into a smaller feature map of 'edge gradients'*

Pooling is optionally applied

Convolution Neural Network (CNN)

In CNNs the filter values are weight parameters that are learned (**feature discovery**)

W_{11}	W_{12}	W_{13}
W_{21}	W_{22}	W_{23}
W_{31}	W_{32}	W_{33}

Convolution Neural Network (CNN)

In CNNs the filter values are weight parameters that are learned (**feature discovery**)

W_{11}	W_{12}	W_{13}
W_{21}	W_{22}	W_{23}
W_{31}	W_{32}	W_{33}

A convolution layer is a set of feature maps, where each map is derived from convolution of 1 filter with input

Convolution Neural Network (CNN)

More hyperparameters:

Size of filter (smaller is more general)

Convolution Neural Network (CNN)

More hyperparameters:

- Size of filter (smaller, like 3x3, is more general)

- Number of pixels to slide over (1 or 2 is usually fine)

Convolution Neural Network (CNN)

More hyperparameters:

- Size of filter (smaller, like 3x3, is more general)

- Number of pixels to slide over (1 or 2 is usually fine)


- Max pooling or not (usually some pooling layers)

- Number of filters (depends on the problem!)

- **A large CNN example**


Convolution with image

- Make 1 layer, using HxWx3 image (3 for Red,Green,Blue channels)




Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)




Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)




Convolution with image

- Make 1 layer, using HxWx3 image (3 for RGB channels)




Convolution with image

- Make 1 layer, using $H \times W \times 3$ image (3 for RGB channels)




Large Scale Versions

- Large Convolution Networks – Alexnet, VGG19, ResNet, GoogLeNet, ...



First convolution layer filters are simple features




What Learned
Convolutions
Look Like

What Learned Convolutions

First convolution layer filters are simple features



Higher layers are more abstract features (or feature combinations)



Convolution Neural Network Summary

CNNs work because convolution layers have a special architecture and function – it is biased to do certain kind of transformations

Low layers have less filters that represent simple local features for all classes

Higher layers have more filters that cover large regions that represent object class features

What is deep learning?

Deep learning refers to learning complex and varied transformations of the input

Deep learning refers to **discovering** useful features of the input

Deep learning is a neural network with many layers

Where to go from here

- Find relevant examples to your domain or task
- Tensorflow has many examples with tutorials in their documentation
- Tensorflow hub and model examples have code and pretrained models

https://tfhub.dev/google/imagenet/inception_v1/classification/4

<https://keras.io/examples/>

- **Next, notebook exercise**

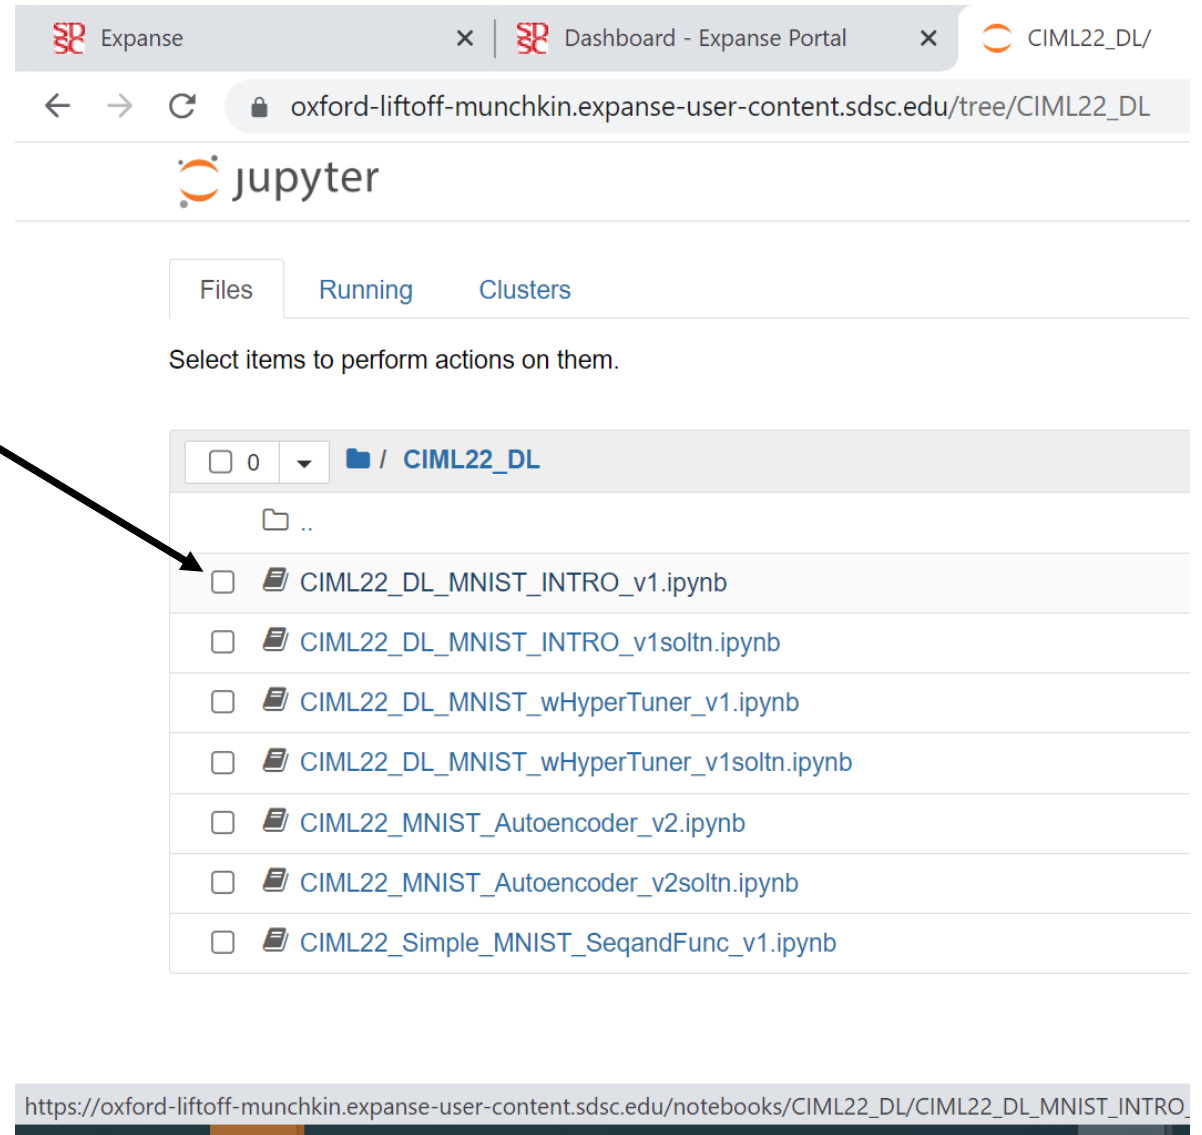
Exercise CNN for Digit Classification

- The ‘hello world’ of CNNs
- It uses MNIST dataset and Keras/Tensorflow
- Goal: Get familiar with Keras and CNN layers coding, and CNN solutions
- We will login and start a notebook (see next pages for quick overview)



In jupyter notebook session
open the MNIST_Intro
notebook

Follow instructions in the
notebook



The screenshot shows a web browser with three tabs: 'Expanse', 'Dashboard - Expanse Portal', and 'CI ML22_DL/'. The address bar shows the URL 'oxford-liftoff-munchkin.expanse-user-content.sdsc.edu/tree/CI ML22_DL'. The JupyterLab interface has tabs for 'Files', 'Running', and 'Clusters'. Below these tabs is the text 'Select items to perform actions on them.' The file browser shows a directory structure with a folder icon and '..' at the top, followed by a list of files: 'CI ML22_DL_MNIST_INTRO_v1.ipynb', 'CI ML22_DL_MNIST_INTRO_v1soltn.ipynb', 'CI ML22_DL_MNIST_wHyperTuner_v1.ipynb', 'CI ML22_DL_MNIST_wHyperTuner_v1soltn.ipynb', 'CI ML22_MNIST_Autoencoder_v2.ipynb', 'CI ML22_MNIST_Autoencoder_v2soltn.ipynb', and 'CI ML22_Simple_MNIST_SeqandFunc_v1.ipynb'. An arrow points from the text 'open the MNIST_Intro notebook' to the first file in the list.

https://oxford-liftoff-munchkin.expanse-user-content.sdsc.edu/notebooks/CI ML22_DL/CI ML22_DL_MNIST_INTRO_v1.ipynb

Keras code for a convolution neural network

```
# -----Set up Model -----
def build_model(numfilters):
    mymodel = keras.models.Sequential()
    mymodel.add(keras.layers.Convolution2D(numfilters,      #<<<< ----- 1
                                           (3, 3),
                                           strides=1,
                                           data_format="channels_last",
                                           activation='relu',
                                           input_shape=(28,28,1)))

    #add another conv layer?  mymodel.add(keras.layers.Convolution2D( ...

    mymodel.add(keras.layers.MaxPooling2D(pool_size=(2,2),strides=2,data_format="channels_la
    mymodel.add(keras.layers.Flatten())                #reorganize 2DxFilters output into 1D

    #-----Now add final classification layers
    mymodel.add(keras.layers.Dense(32, activation='relu'))
    mymodel.add(keras.layers.Dense(10, activation='softmax'))

    # ----- Now configure model algorithm -----
    mymodel.compile(loss='categorical_crossentropy',
                    optimizer=keras.optimizers.Adam(learning_rate=0.001),
```

A sequential model

Add convolution layer

*Add max pooling, then
flatten into a vector for
classification layers*

- Remember every layer has some input, output
- Keras figures out the shapes
- Not every layer in Keras has trainable parameters – like which ones above?

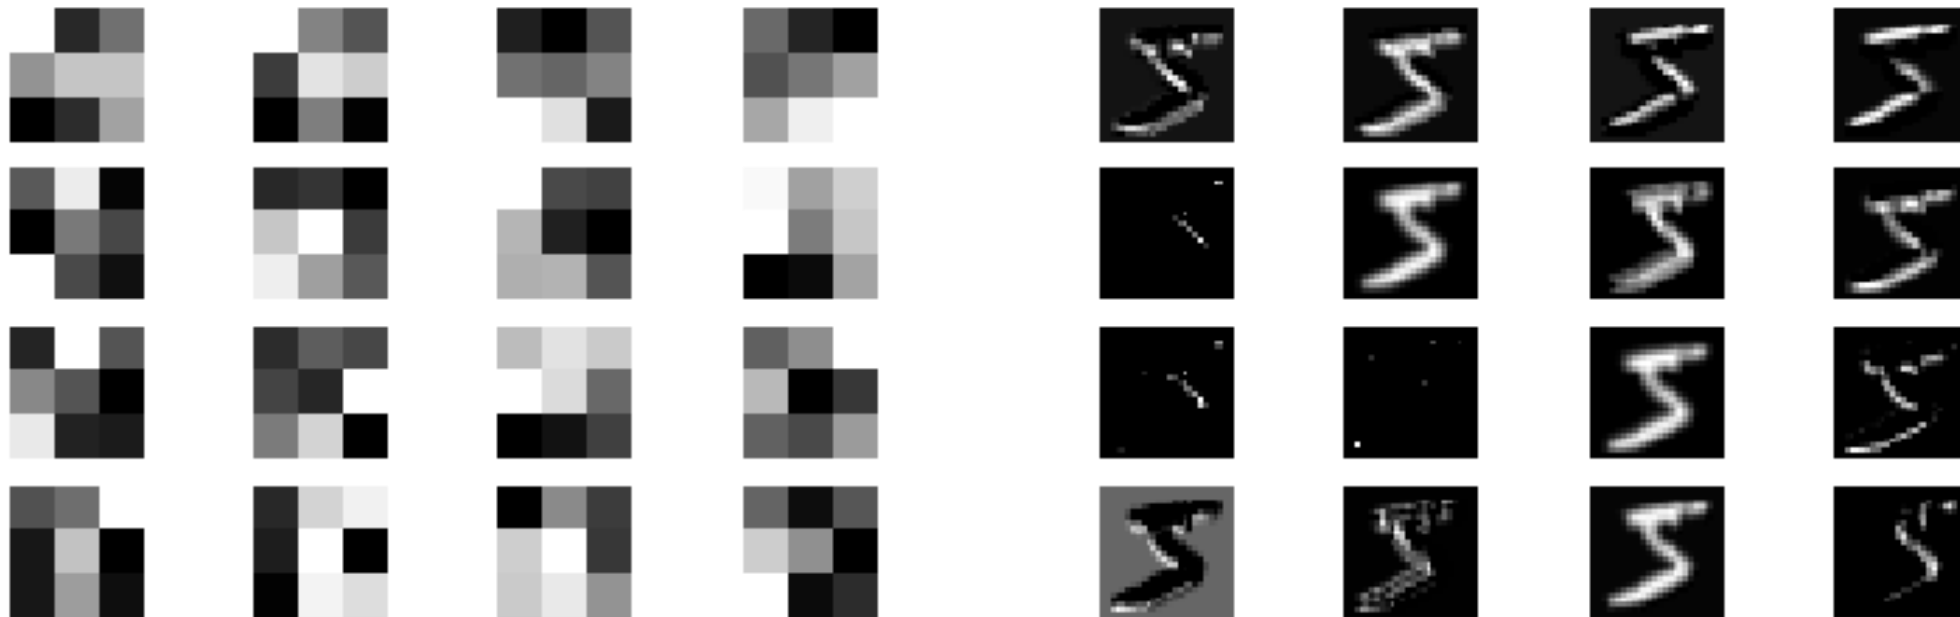
Zooming in on keras convolution layers statements

Use 16 filters, each of size 3x3

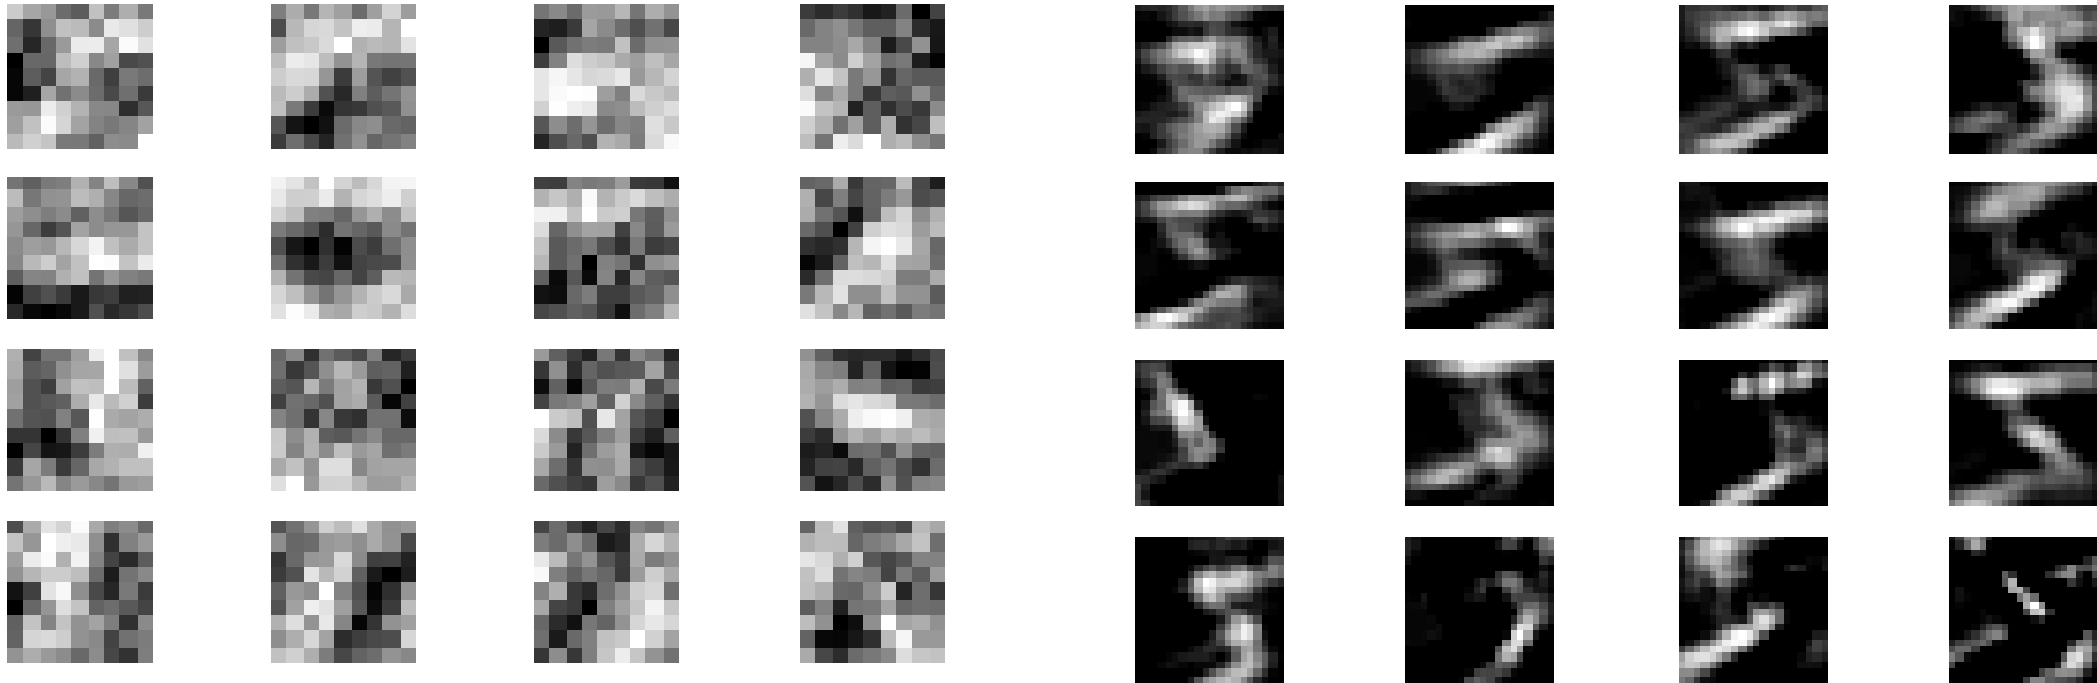
```
my_model.add(tf.keras.layers.Convolution2D(filters=16,  
                                            kernel_size=(3, 3),  
                                            strides=1,  
                                            data_format="channels_last",  
                                            activation='relu',  
                                            input_shape=(28,28,1)))
```

*Input shape does not
include number of images*

Exercise notes: 3x3 first convolution layer filter and activation



9x9 first convolution layer filter and activation



End